

Task 1 - Language Benchmark of matrix  
multiplication  
Big Data

Alberto Guedes

October 2025

## Contents

<b>1</b>	<b>Methodology</b>	<b>4</b>
1.1	Implementations . . . . .	4
1.2	Benchmark harness . . . . .	4
<b>2</b>	<b>Results</b>	<b>4</b>
<b>3</b>	<b>Discussion</b>	<b>7</b>
3.1	Runtime . . . . .	7
3.2	CPU and memory . . . . .	7
3.3	Threats to validity . . . . .	7
<b>4</b>	<b>Conclusion</b>	<b>7</b>

**Abstract**

The following report presents the development and execution of benchmarking tests for matrix multiplication implemented in three different programming languages: C, Java, and Python. The main objective is to compare their performance in terms of execution time and computational efficiency when performing large-scale numerical operations. The results obtained provide insights into the trade-offs between language design, memory management, and runtime performance, offering a deeper understanding of how each language handles intensive computational workloads. The source code can be found in the following repository: [GitHub Repository](#)

# 1 Methodology

## 1.1 Implementations

The three implementations compute  $C \leftarrow A \times B$  using the classical triple loop.

- C: static arrays allocated as 2D blocks. Time measured with a high-resolution monotonic clock.
- Java: arrays of `double[]` with warm-up on the first run. Built with Java 21.
- Python: lists-of-lists using pure Python loops. Time measured with `time.perf_counter()`.

## 1.2 Benchmark harness

For each language and matrix size  $n \in \{10, 50, 100, 256, 512\}$ , 10 repetitions have been run, printing a CSV row per run. Wall time is the primary metric. CPU and memory are also emitted:

- Python/Java: CPU as percent (%), memory as MB (RSS-like).
- C: CPU reported as accumulated CPU seconds, memory in MB.

CSV artifacts are stored under `results/`. A companion notebook including data visualization (`Benchmark_results.ipynb`).

# 2 Results

Table 1 shows mean wall time (seconds) for each language and size (10 runs per size). C is the fastest at all sizes; Java trails C by a small constant factor; Python exhibits orders of magnitude slower runtime as  $n$  grows, consistent with interpreter overhead.

Table 1: Mean wall time (seconds) by language and matrix size.

oprule Language	$n = 10$	$n = 50$	$n = 100$	$n = 256$	$n = 512$
C	0.000002	0.000108	0.001332	0.021497	0.180225
Java	0.000804	0.000923	0.002767	0.025719	0.285948
Python	0.000102	0.011126	0.088605	1.542429	14.234368

CPU and memory metrics follow in Tables 2 and 3. Note that C reports CPU in seconds, whereas Python/Java report percent; the values are therefore not directly comparable across languages. Within-language trends still inform relative scaling.

Table 2: Mean CPU metric by language and size. Python/Java in percent (%), C in CPU seconds.

oprule Language	$n = 10$	$n = 50$	$n = 100$	$n = 256$	$n = 512$
C (s)	0.00017	0.01074	0.13233	2.14209	17.82012
Java (%)	1.00	11.50	27.78	15.14	12.35
Python (%)	5.00	12.94	14.97	18.71	17.23

Table 3: Mean memory usage (MB) by language and size.

oprule Language	$n = 10$	$n = 50$	$n = 100$	$n = 256$	$n = 512$
C	0.00	0.00	0.00	0.00	0.00
Java	0.09	0.02	0.05	0.09	0.05
Python	0.00	0.01	0.03	0.24	4.74

Figure 1 (produced by `Benchmark_results.ipynb`) visualizes mean time, CPU, and memory versus size on log-scaled axes for readability.




Figure 1: Benchmark results for C, Java and Python implementations.

## 3 Discussion

### 3.1 Runtime

The naive C implementation consistently outperforms Java and Python due to ahead-of-time compilation and minimal overhead. Java is competitive, particularly at small to mid sizes, but remains slower than C. Python’s interpreter overhead dominates at larger  $n$ , yielding multi-order-of-magnitude slowdowns.

### 3.2 CPU and memory

Java and Python show modest CPU percentages per run, reflecting short single-threaded bursts per measurement. C’s CPU metric (seconds) increases roughly with  $\mathcal{O}(n^3)$  work, tracking wall time. Python memory grows with  $n$  as expected from list-of-list allocations; Java’s reported per-run memory increments remain small given JVM reuse and GC behavior during runs; C’s per-run deltas are effectively negligible at this scale.

### 3.3 Threats to validity

Results depend on compiler flags, JVM warm-up, Python interpreter version, and system load. CPU metrics are not normalized across languages (seconds vs percent) and should not be compared directly. Cache effects and memory layout (row-major access patterns) may bias certain loop orders. The implementations are intentionally naive and not vectorized nor parallelized.

## 4 Conclusion

C delivers the best performance for naive matrix multiplication, Java follows at a small constant factor, and Python is suitable only for small sizes unless vectorized libraries (NumPy) are used. For production workloads, prefer optimized libraries or low-level languages; for pedagogy or prototyping, Python remains convenient but slow.

## Reproducibility

Benchmarks were run with 10 repetitions per size using the scripts and programs in this repository. Raw CSVs are under `results/`. The notebook `Benchmark_results.ipynb` computes aggregates and produces comparative plots.